

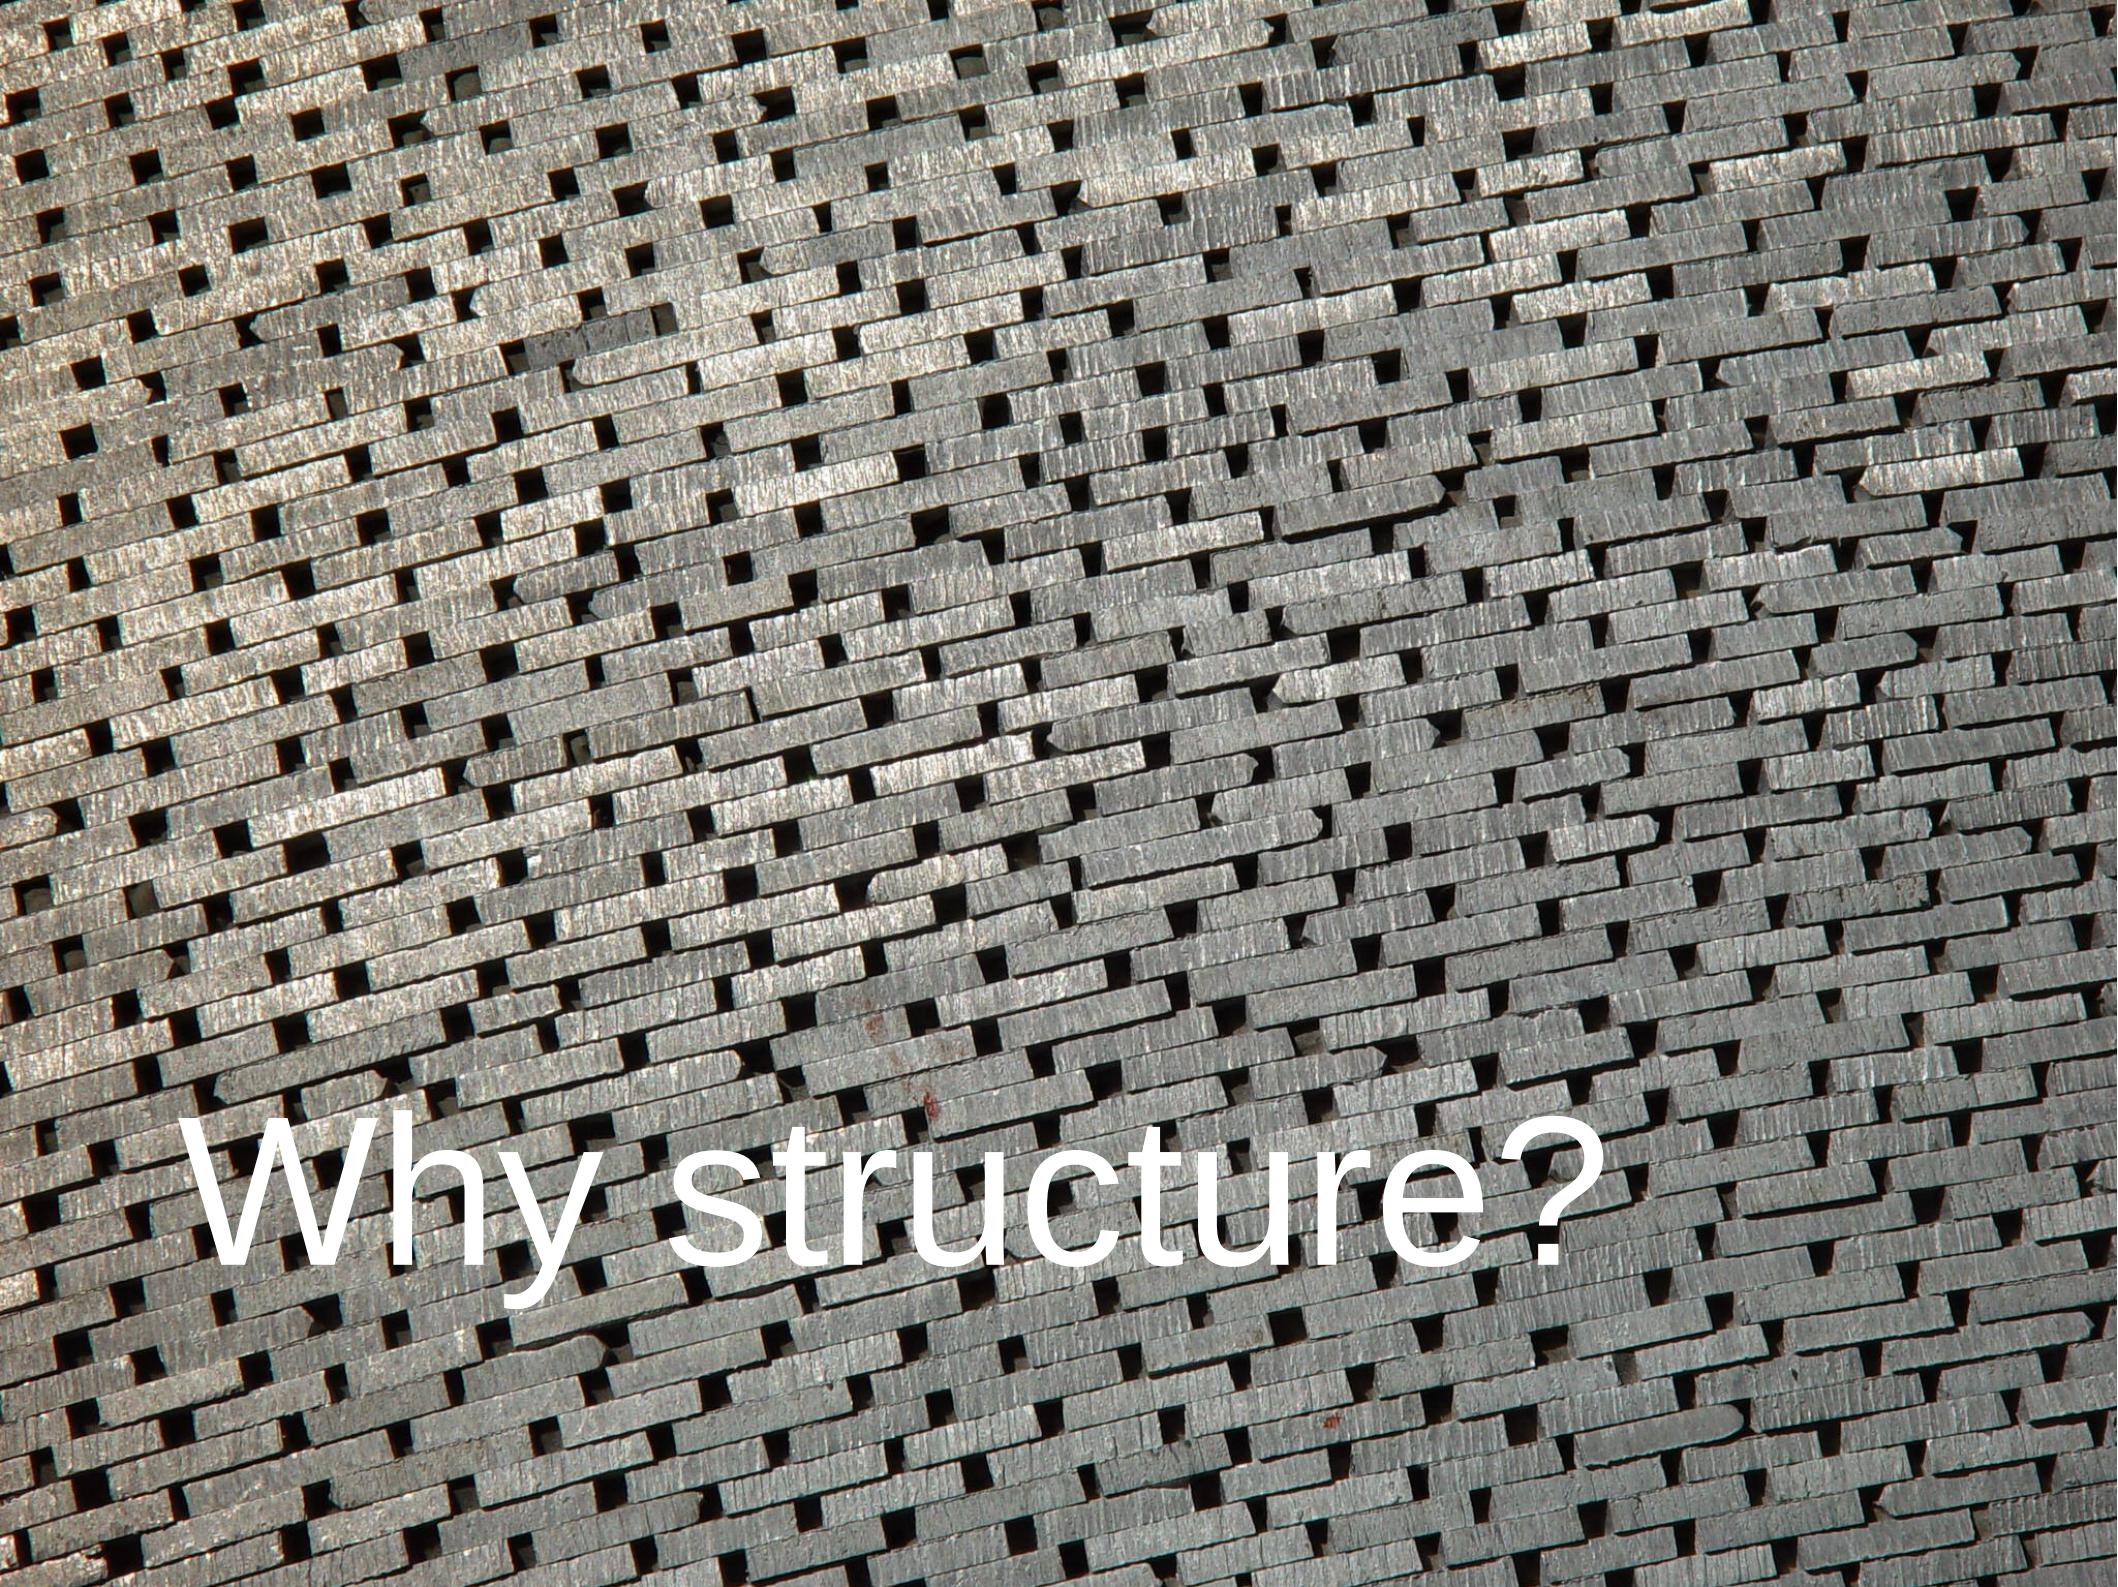
PyStruct

Structured Prediction in Python

Andreas Mueller (NYU Center for Data Science, scikit-learn)

Structured Prediction

$$y = (y_1, y_2, \dots y_{n_k})$$



Why structure?

Applications: Multi-Label Classification

	Politics	Sports	Finance	Domestic	Religion
News Story1	1	0	0	1	1
News Story2	0	1	0	1	0
News Story3	0	0	1	0	0

Applications: Multi-Label Classification

	Politics	Sports	Finance	Domestic	Religion
News Story1	1	0	0	1	1
News Story2	0	1	0	1	0
News Story3	0	0	1	0	0

	Owns Car	Smokes	Married	Self-Employed	Has Kids
Customer1	1	0	1	0	1
Customer2	1	1	0	1	0
Customer3	0	1	1	0	0

Applications: Sequence Tagging



Applications: Sequence Tagging



Stroke cat.



Stroke cat.



Stroke cat.

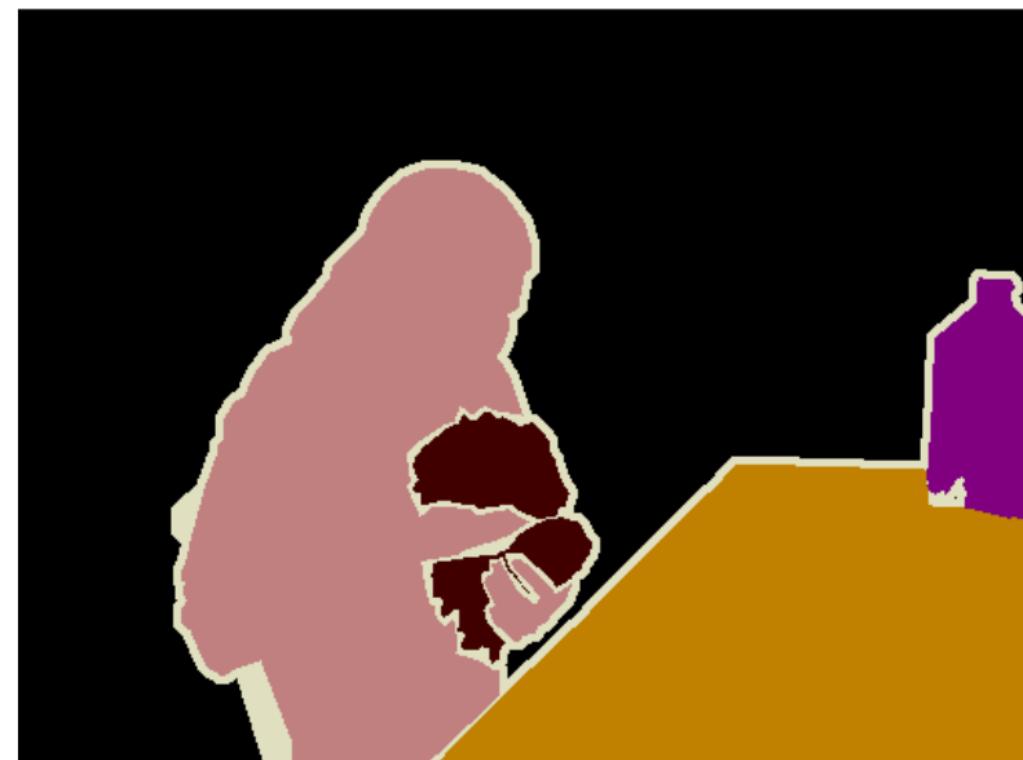


Open trash can.



Put cat in trash can.

Applications: Image Segmentation





CAUTION: MATH

The Essence of Structured Prediction

$$f(x, w) := \arg \max_{y \in \mathcal{Y}} g(x, y, w)$$

The Essence of Structured Prediction

$$f(x, w) := \arg \max_{y \in \mathcal{Y}} g(x, y, w)$$

If you like:

$$\arg \max_{y \in \mathcal{Y}} p(y|x, w)$$

Pairwise Structured Models

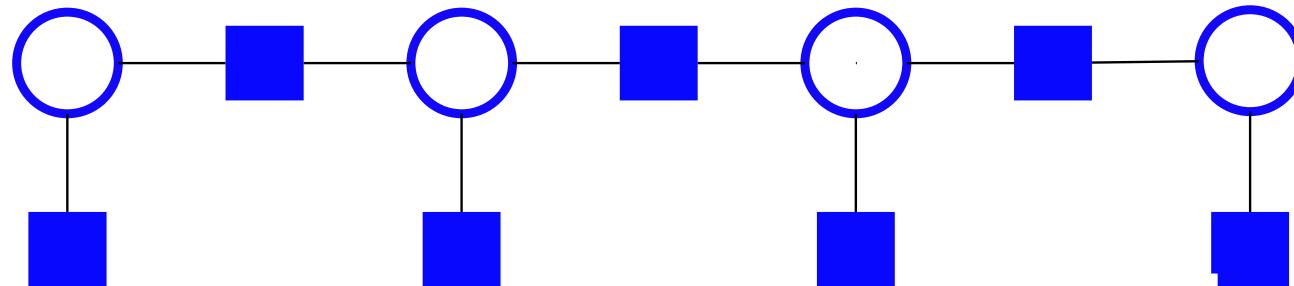
$$\arg \max_{y_1, y_2, \dots, y_n} w^T \psi(x, y)$$

$$= \arg \max_{y_1, y_2, \dots, y_n} \sum_I w_i^T \psi(x, y_i) + \sum_{(i,j) \in E} w_{i,j}^T \psi(x, y_i, y_j)$$

Pairwise Structured Models

$$\arg \max_{y_1, y_2, \dots, y_n} w^T \psi(x, y)$$

$$= \arg \max_{y_1, y_2, \dots, y_n} \sum_I w_i^T \psi(x, y_i) + \sum_{(i,j) \in E} w_{i,j}^T \psi(x, y_i, y_j)$$



PyStruct Architecture

Estimator = Learner + Model + Inference

$$\arg \max_{y_1, y_2, \dots, y_n} w^T \psi(x, y)$$

PyStruct Architecture

Estimator = Learner + Model + Inference

$$\arg \max_{y_1, y_2, \dots, y_n} w^T \psi(x, y)$$

```
model = ChainCRF(inference="max_product")
ssvm = OneSlackSSVM(model=model, C=.1, inference_cache=50,
                     tol=0.1, verbose=3)
ssvm.fit(X_train, y_train)
```

Implemented Methods

Estimator = Learner + Model + Inference

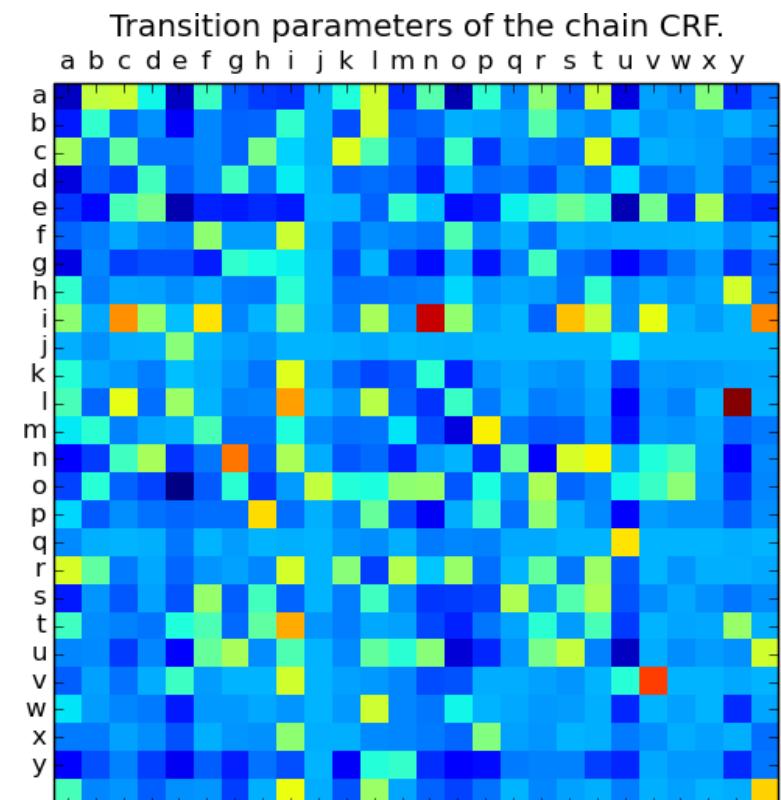
- Learner: SubgradientSSVM, StructuredPerceptron, OneSlackSSVM, LatentSSVM
- Model: BinaryClf, MultiLabelClf, ChainCRF, GraphCRF, EdgeFeatureGraphCRF
- Inference: Linear Programming, QPBO (PyQPBO), Dual Decomposition (AD3), Message Passing, everything (OpenGM)

Sequence Tagging example



```
model = ChainCRF(inference="max_product")
ssvm = OneSlackSSVM(model=model, C=.1, inference_cache=50,
                     tol=0.1, verbose=3)
ssvm.fit(X_train, y_train)
```

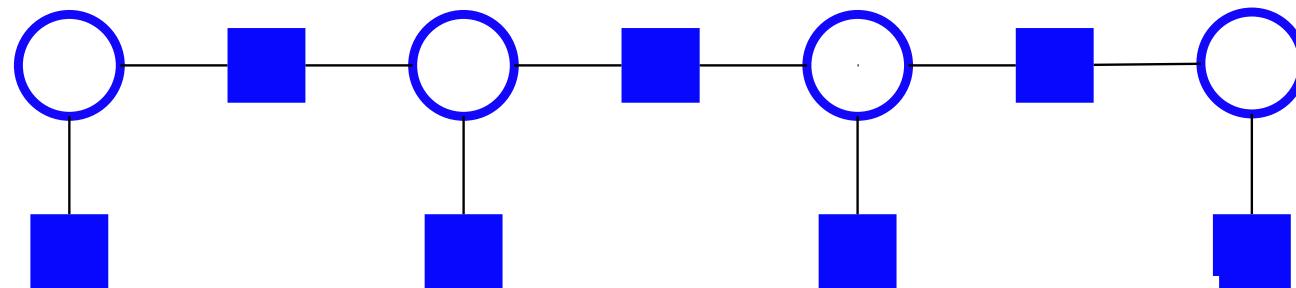
Sequence Tagging example



```
model = ChainCRF(inference="max_product")
ssvm = OneSlackSSVM(model=model, C=.1, inference_cache=50,
                     tol=0.1, verbose=3)
ssvm.fit(X_train, y_train)
```

The Devil is in the Inference

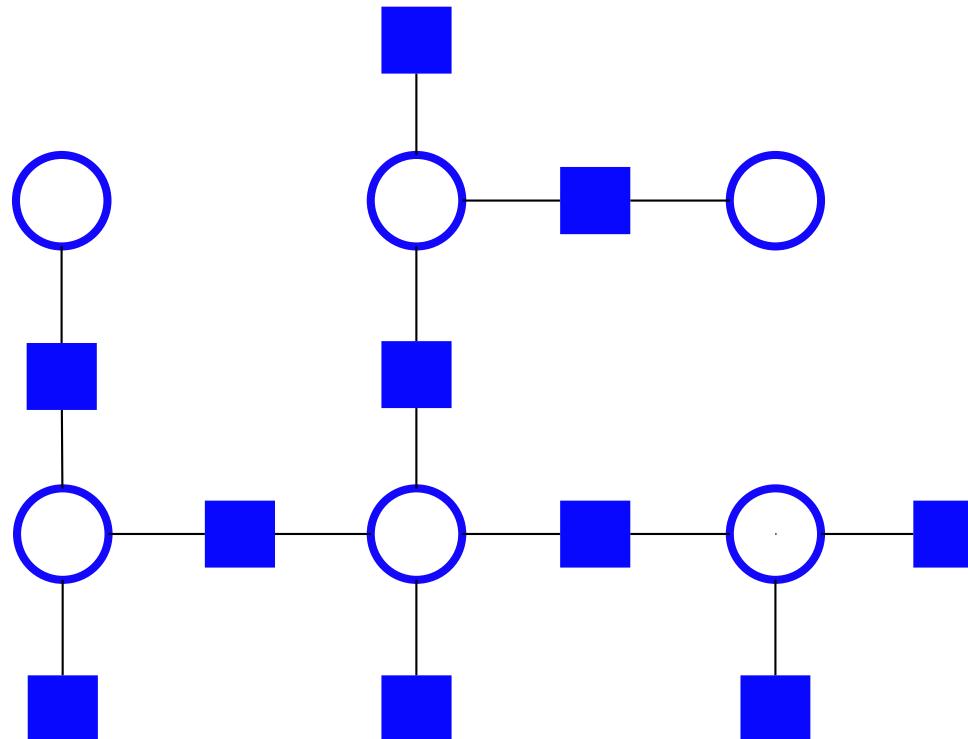
$$\arg \max_{y_1, y_2, \dots, y_n} w^T \psi(x, y)$$



Easy: Dynamic Programming

The Devil is in the Inference

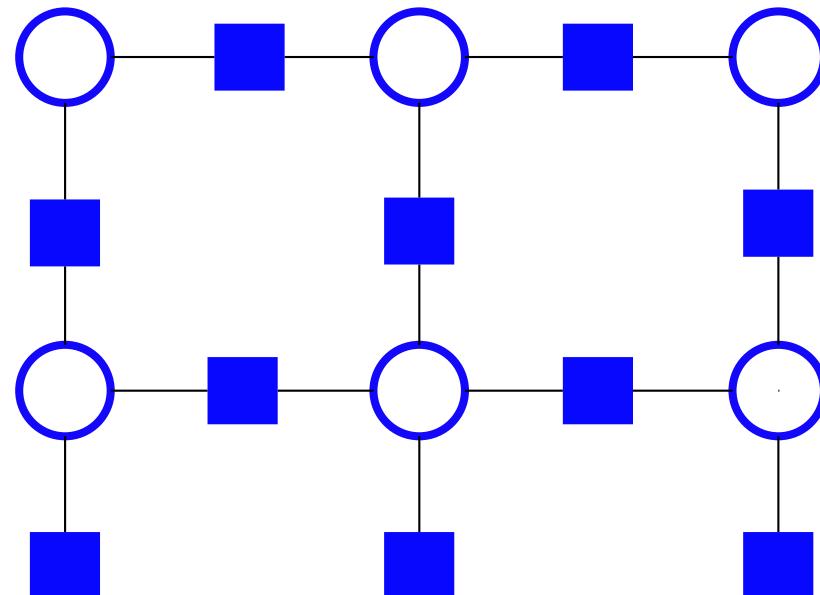
$$\arg \max_{y_1, y_2, \dots, y_n} w^T \psi(x, y)$$



Easy: Dynamic Programming

The Devil is in the Inference

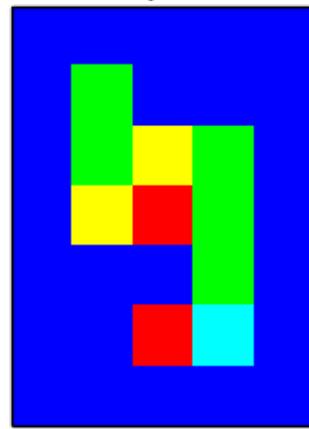
$$\arg \max_{y_1, y_2, \dots, y_n} w^T \psi(x, y)$$



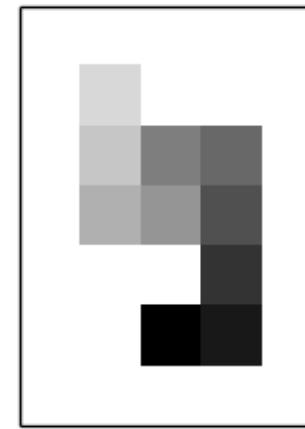
HARD!
AD3, QPBO, LP, Loopy BP,

Grid Graphs: Snakes

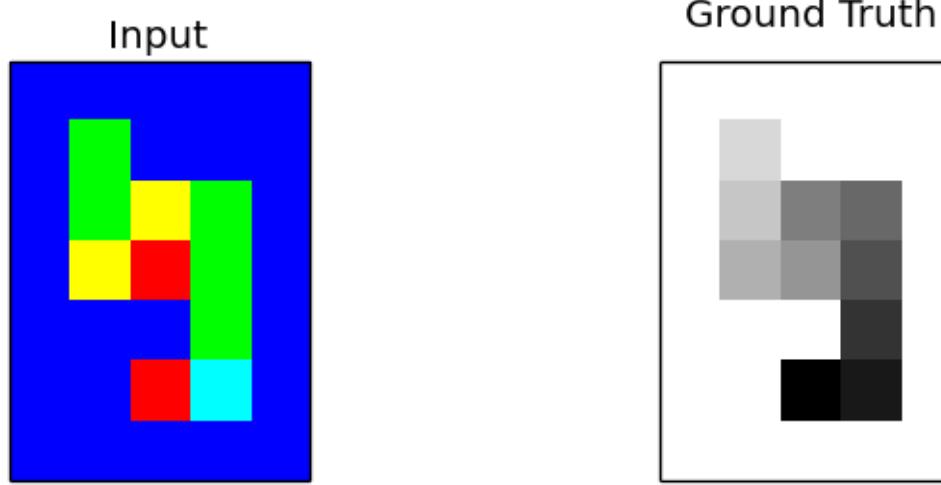
Input



Ground Truth



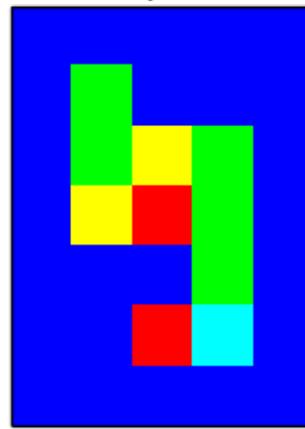
Grid Graphs: Snakes



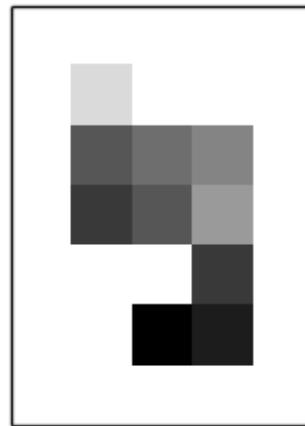
```
crf = EdgeFeatureGraphCRF(inference_method='qpbo')
ssvm = OneSlackSSVM(crf, inference_cache=50, C=.1, tol=.1,
                     switch_to='ad3', n_jobs=1)
ssvm.fit(X_train_edge_features, Y_train_flat)
```

Grid Graphs: Snakes

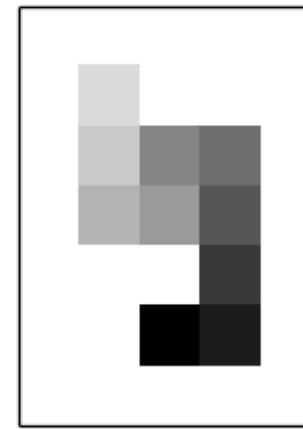
Input



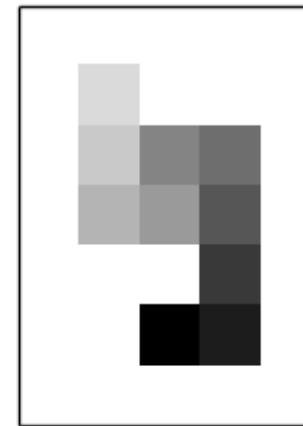
Prediction w/o edge features



Ground Truth



Prediction with edge features



Classes of Inference Algorithms

Exact Algorithms

Max-Product (Chains, Trees) 'max-product'

Exhaustive (usually too expensive)

Relaxed algorithms + branch & bound

('ad3', {'branch_and_bound': True})

Relaxed

Linear Programming (slow) 'lp'

Dual Decomposition 'ad3'

Approximate / heuristics

Loopy message passing 'max-product'

QPBO 'qpbo'

Classes of Inference Algorithms

Exact Algorithms

Max-Product (Chains, Trees) 'max-product'

Exhaustive (usually too expensive)

Relaxed algorithms + branch & bound

('ad3', {'branch_and_bound': True})

Relaxed

Linear Programming (slow) 'lp'

Dual Decomposition 'ad3'

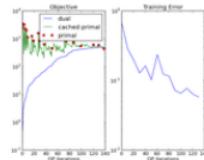
Approximate / heuristics

Loopy message passing 'max-product'

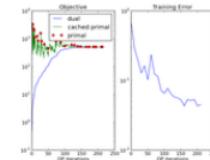
QPBO 'qpbo'

Install OpenGM for many more!

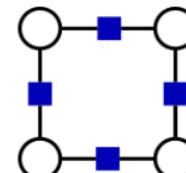
Examples



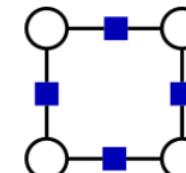
Plotting the objective and constraint caching in 1-slab SSVMs



Efficient exact learning of 1-slab SSVMs



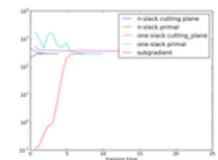
SVM as CRF



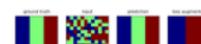
Semantic Image Segmentation on Pascal VOC



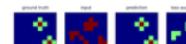
Latent Dynamics CRF



SVM objective values



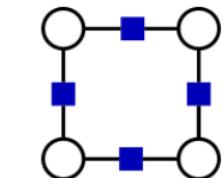
Learning directed interactions on a 2d grid



Learning interactions on a 2d grid



Latent SVM for odd vs. even digit classification



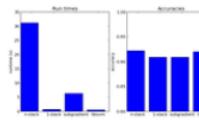
Mult-label classification



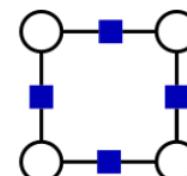
Latent Variable Hierarchical CRF



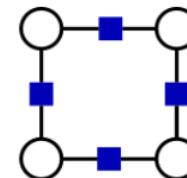
OCR Letter sequence recognition



Binary SVM as SSVM



Crammer-Singer Multi-Class SVM



Comparing PyStruct and SVM-Struct

Thank you for your attention.



@t3kcit



@amueller



t3kcit@gmail.com



<http://amueller.github.io>